

Constraint Satisfaction Problems 2

CS 4300
Artificial Intelligence

Prof. Alan Kuntz
alan.kuntz@utah.edu

Announcements

- P2 due Friday 10/6 BUT
 - The entirety of Fall break will count as 1 penalty free late day.

Ordering: Minimum Remaining Values


- Variable Ordering: Minimum remaining values (MRV):
 - Choose the variable with the fewest legal left values in its domain



- Why min rather than max?
- Also called “most constrained variable”
- “Fail-fast” ordering



Ordering: Least Constraining Value

- Value Ordering: Least Constraining Value
 - Given a choice of variable, choose the *least constraining value*
 - I.e., the one that rules out the fewest values in the remaining variables
 - Note that it may take some computation to determine this! (E.g., rerunning filtering)
- Why least rather than most?
- Combining these ordering ideas makes 1000 queens feasible




Problem Structure

- **Extreme case: independent subproblems**
 - Example: Tasmania and mainland do not interact
- Independent subproblems are identifiable as connected components of constraint graph
- Suppose a graph of n variables can be broken into subproblems of only c variables:
 - Worst-case solution cost is $O((n/c)(d^c))$, much better than $O(d^n)$!
 - E.g., $n = 80$, $d = 2$, $c = 20$
 - $2^{80} = 4$ billion years at 10 million nodes/sec
 - $(4)(2^{20}) = 0.4$ seconds at 10 million nodes/sec
- **Sad news: Doesn't come up that often**



Tree-Structured CSPs

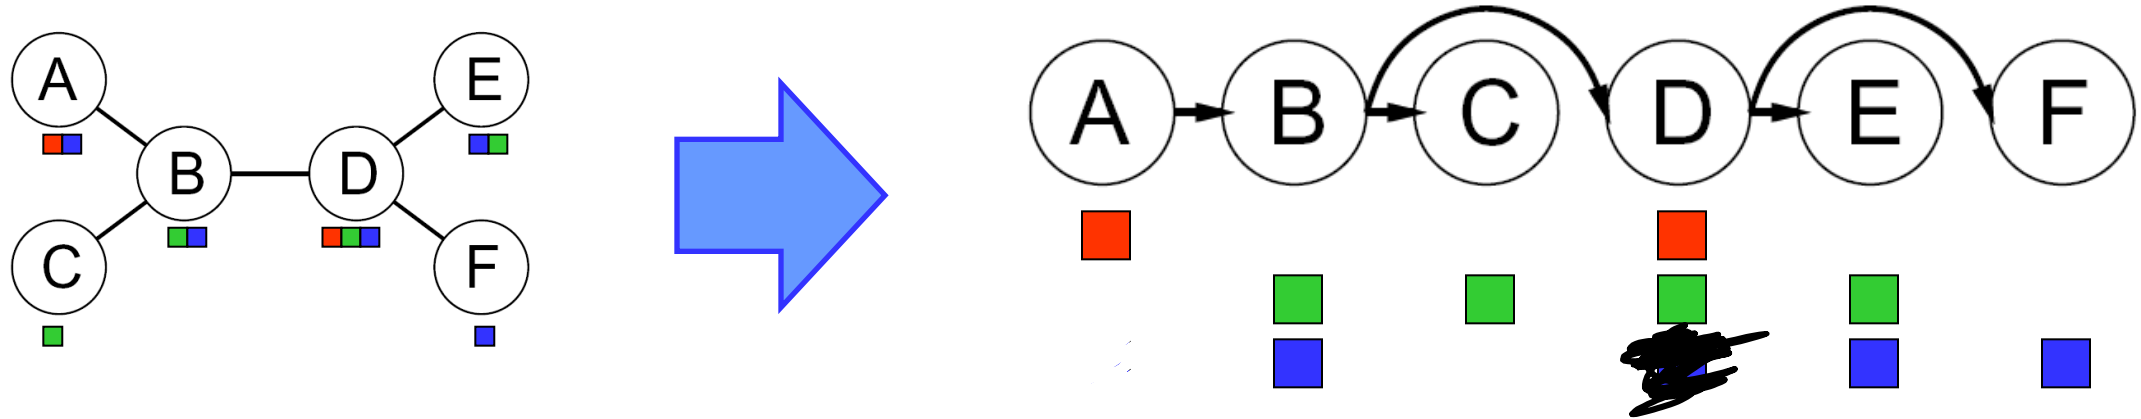


- Theorem: if the constraint graph has no cycles, the CSP can be solved in $O(n d^2)$ time
 - Compare to general CSPs, where worst-case time is $O(d^n)$
- This property also applies to probabilistic reasoning (later)

Tree-Structured CSPs

- Algorithm for tree-structured CSPs:

- Order: Choose a root variable, order variables so that parents precede children




- Remove backward: For $i = n : 2$, apply $\text{RemoveInconsistent}(\text{Parent}(X_i), X_i)$
- Assign forward: For $i = 1 : n$, assign X_i consistently with $\text{Parent}(X_i)$

- Runtime: $O(n d^2)$ (why?)

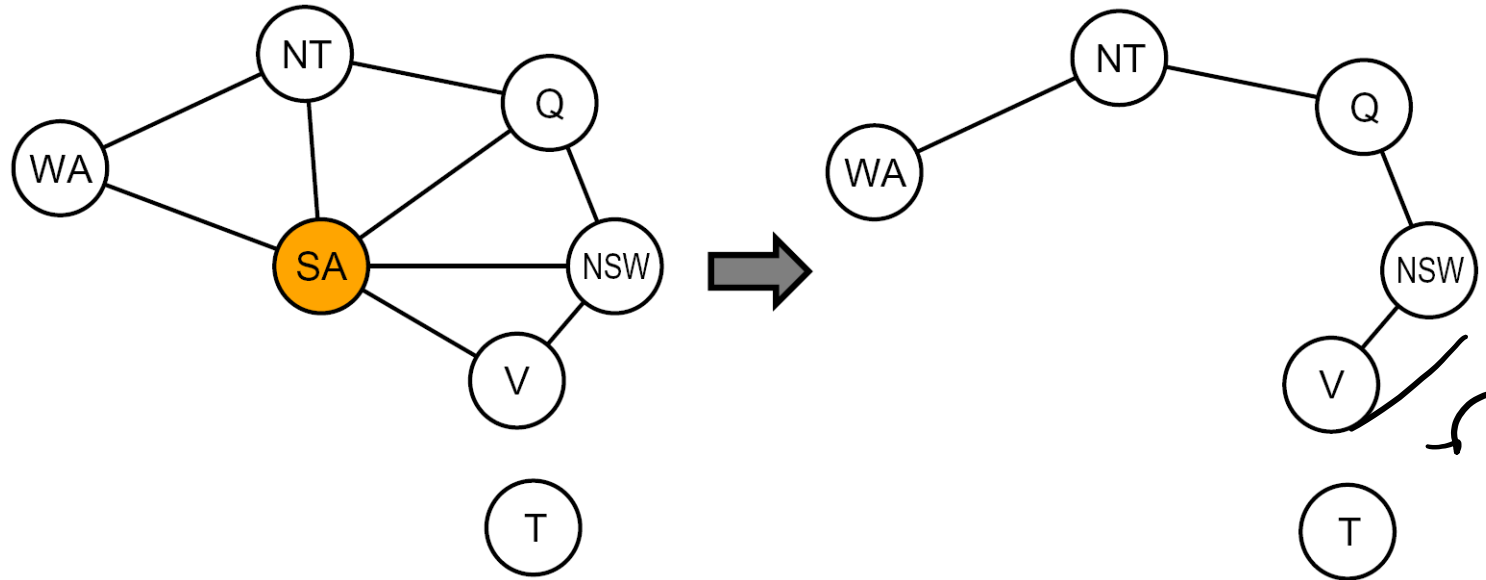
Tree-Structured CSPs

- Claim 1: After backward pass, all root-to-leaf arcs are consistent
- Proof: Each $X \rightarrow Y$ was made consistent at one point and Y 's domain could not have been reduced thereafter (because Y 's children were processed before Y)



- Claim 2: If root-to-leaf arcs are consistent, forward assignment will not backtrack
- Proof: Induction on position
- Why doesn't this algorithm work with cycles in the constraint graph?
- Note: we'll see this basic idea again with Bayes' nets

Nearly Tree-Structured CSPs



- Conditioning: instantiate a variable, prune its neighbors' domains
- Cutset conditioning: instantiate (*in all ways*) a set of variables such that the remaining constraint graph is a tree
- Cutset size c gives runtime $O((d^c)(n-c)d^2)$, very fast for small c


Cutset Conditioning

Choose a cutset

Instantiate the cutset
(all possible ways)


Compute residual CSP
for each assignment

Solve the residual CSPs
(tree structured)



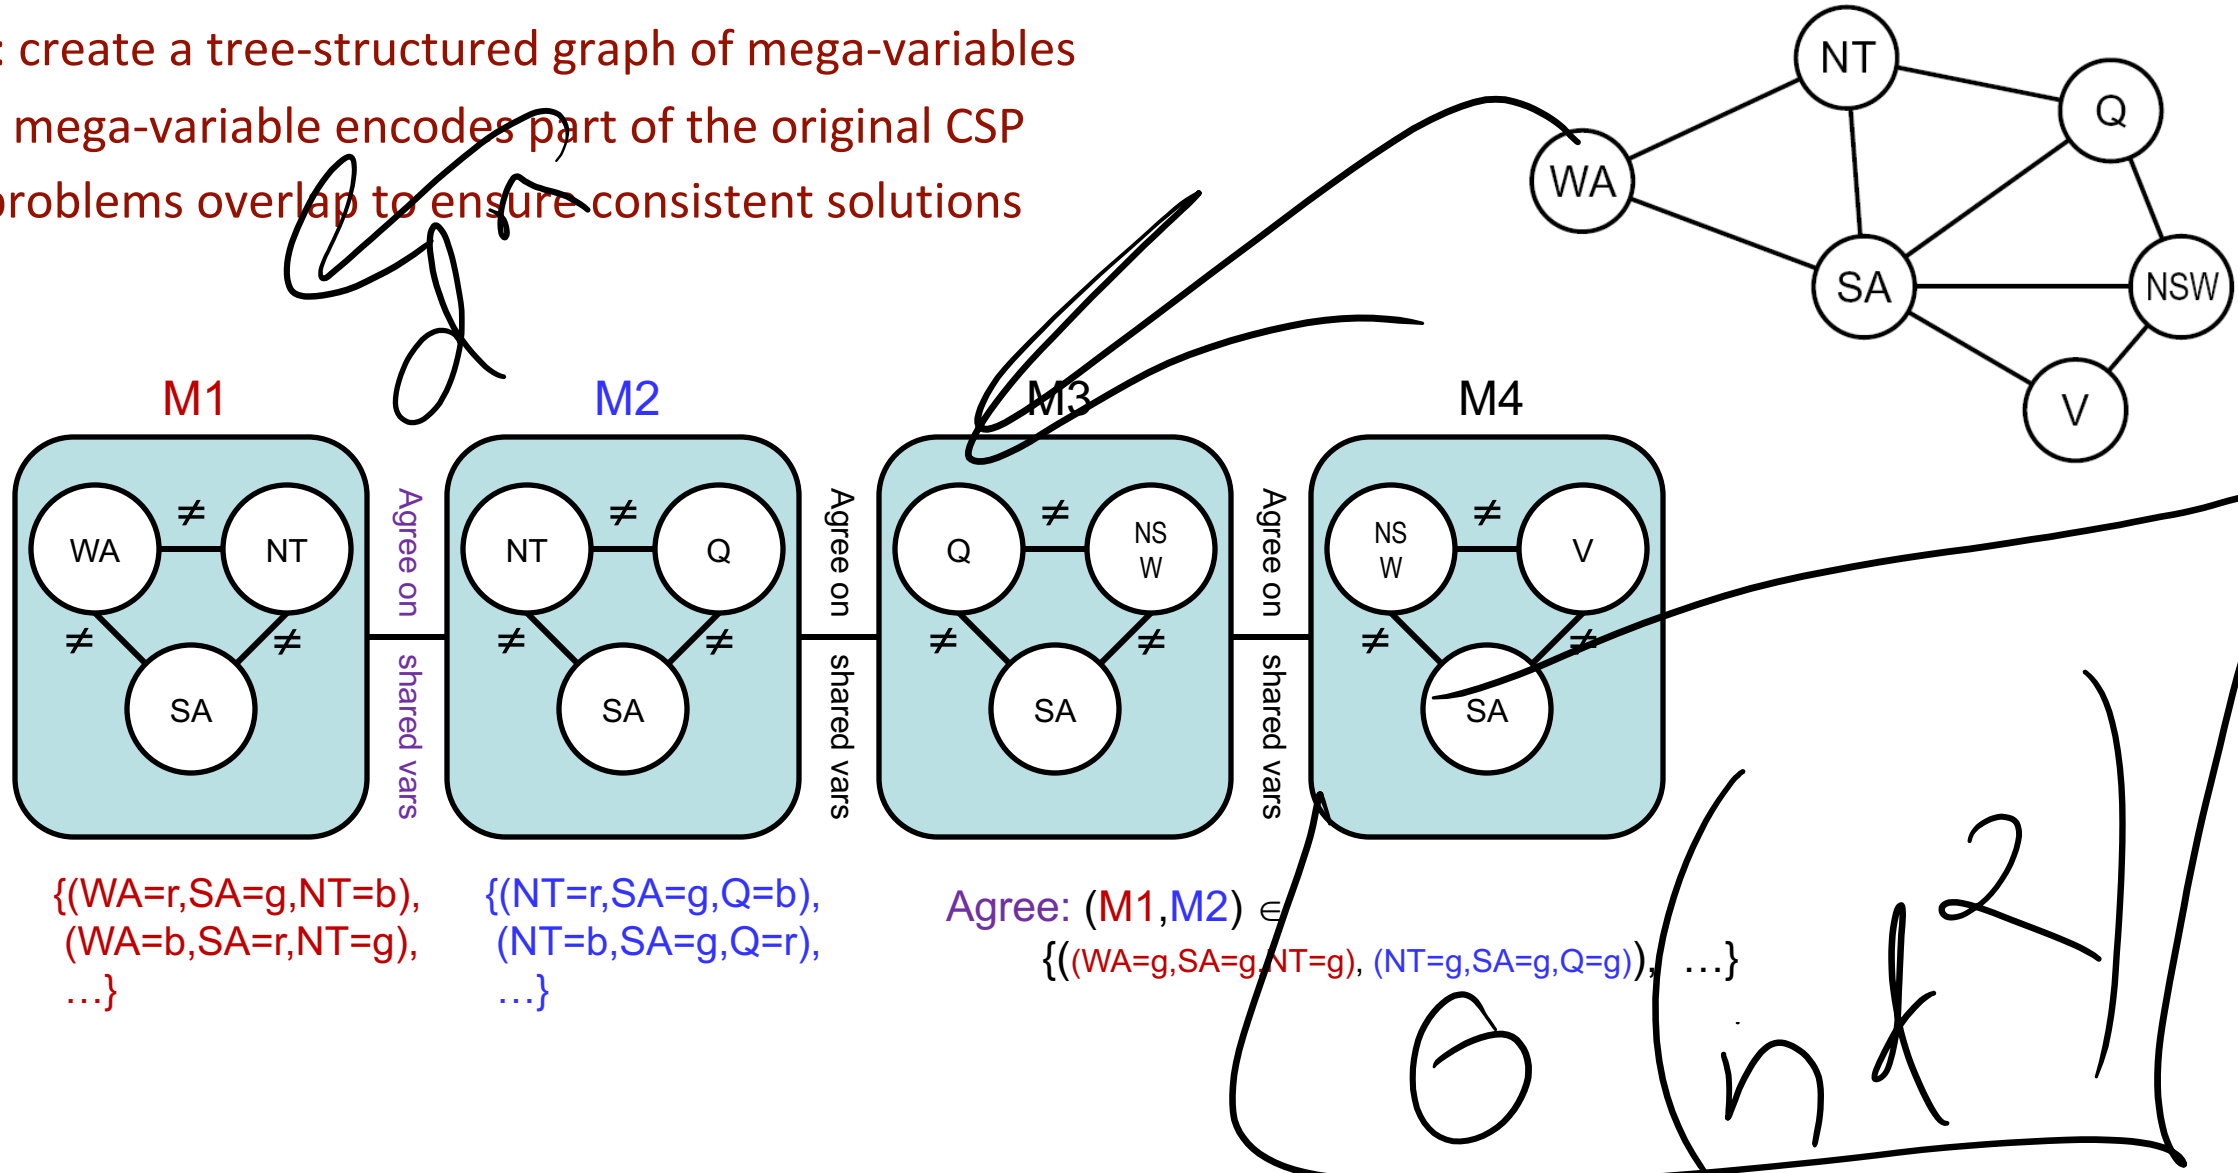
Cutset Quiz

- Find the smallest cutset that results in a tree for the graph below.



Tree Decomposition


- Idea: create a tree-structured graph of mega-variables
- Each mega-variable encodes part of the original CSP
- Subproblems overlap to ensure consistent solutions




Iterative Improvement

Iterative Algorithms for CSPs

- Local search methods typically work with “complete” states, i.e., all variables assigned
- To apply to CSPs:
 - Take an assignment with unsatisfied constraints
 - Operators *reassign* variable values
 - No fringe! Live on the edge.
- Algorithm: While not solved,
 - Variable selection: randomly select any conflicted variable
 - Value selection: min-conflicts heuristic:
 - Choose a value that violates the fewest constraints
 - I.e., hill climb with $h(n)$ = total number of violated constraints




Example: 4-Queens



- States: 4 queens in 4 columns ($4^4 = 256$ states)
- Operators: move queen in column
- Goal test: no attacks
- Evaluation: $c(n)$ = number of attacks

[Demo: coloring – iterative improvement]



Iterative Improvement – Coloring



Performance of Min-Conflicts

- Given random initial state, can solve n-queens in almost constant time for arbitrary n with high probability (e.g., n = 10,000,000)!
- The same appears to be true for any randomly-generated CSP *except* in a narrow range of the ratio

$$R = \frac{\text{number of constraints}}{\text{number of variables}}$$



Summary: CSPs

- CSPs are a special kind of search problem:
 - States are partial assignments
 - Goal test defined by constraints
- Basic solution: backtracking search
- Speed-ups:
 - Ordering
 - Filtering
 - Structure
- Iterative min-conflicts is often effective in practice

Example from my Research

Optimizing Hospital Room Layout to Reduce the Risk of Patient Falls

Sarvenaz Chaeibakhsh¹, Roya Sabbagh Novin¹, Tucker Hermans²,
Andrew Merryweather¹, and Alan Kuntz²


¹*Department of Mechanical Engineering, University of Utah, UT, USA*

²*School of Computing, University of Utah, UT, USA*

sarvenaz.chaeibakhsh@utah.edu


Optimizing Hospital Room Layouts

- Reconfiguring hospital room object placement layout to reduce the risk of fall.
 - Use of **computer layout planning** and **optimization** models to provide a safer room for patient's ambulation

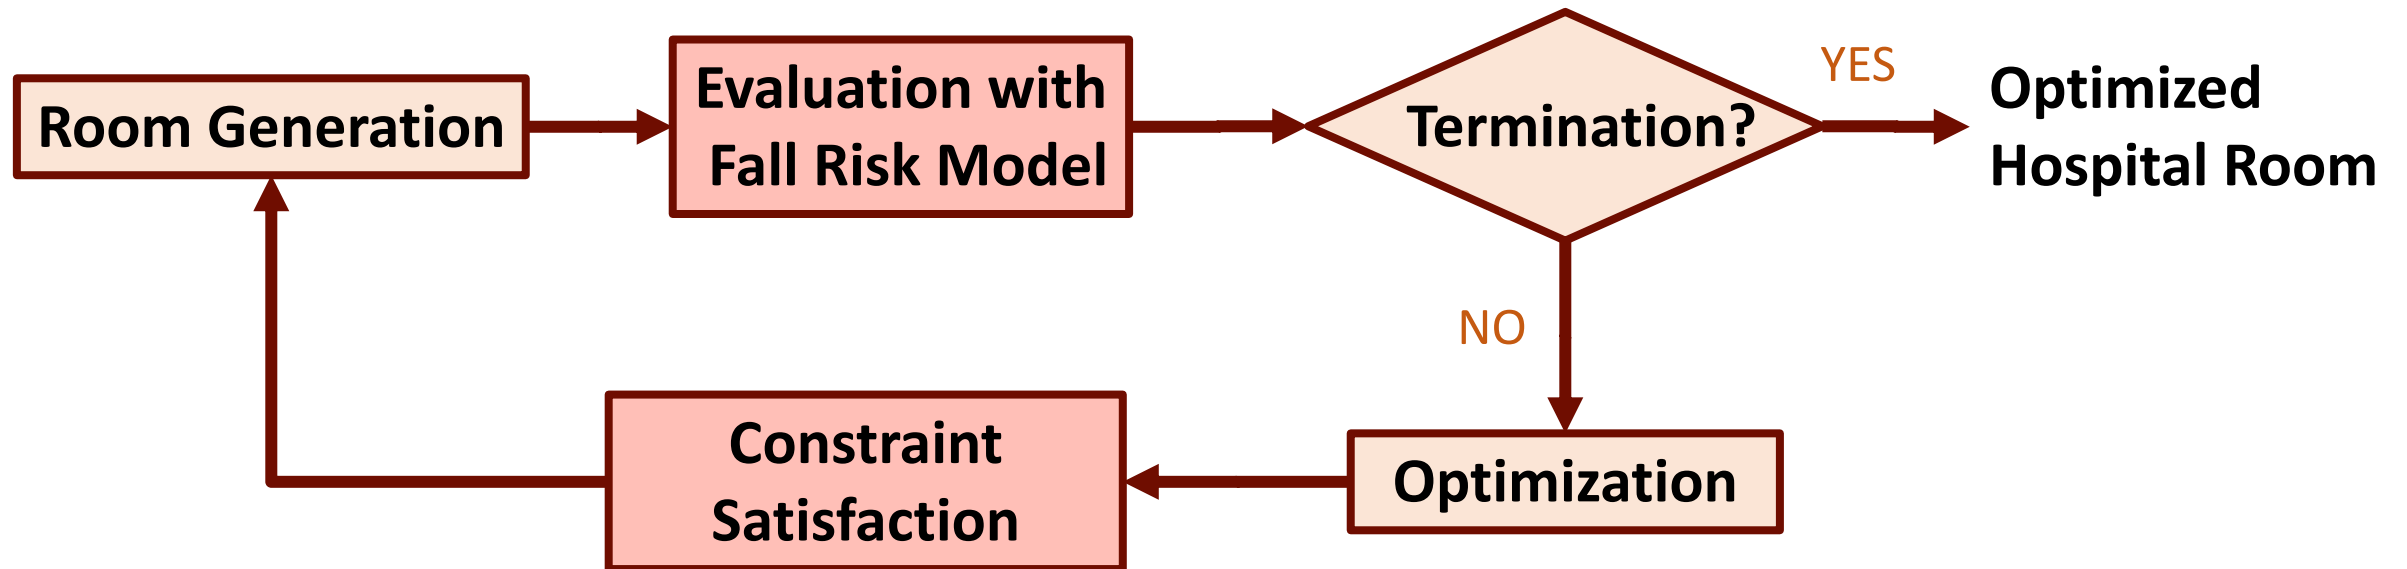


Computer Layout
Planning


Optimization



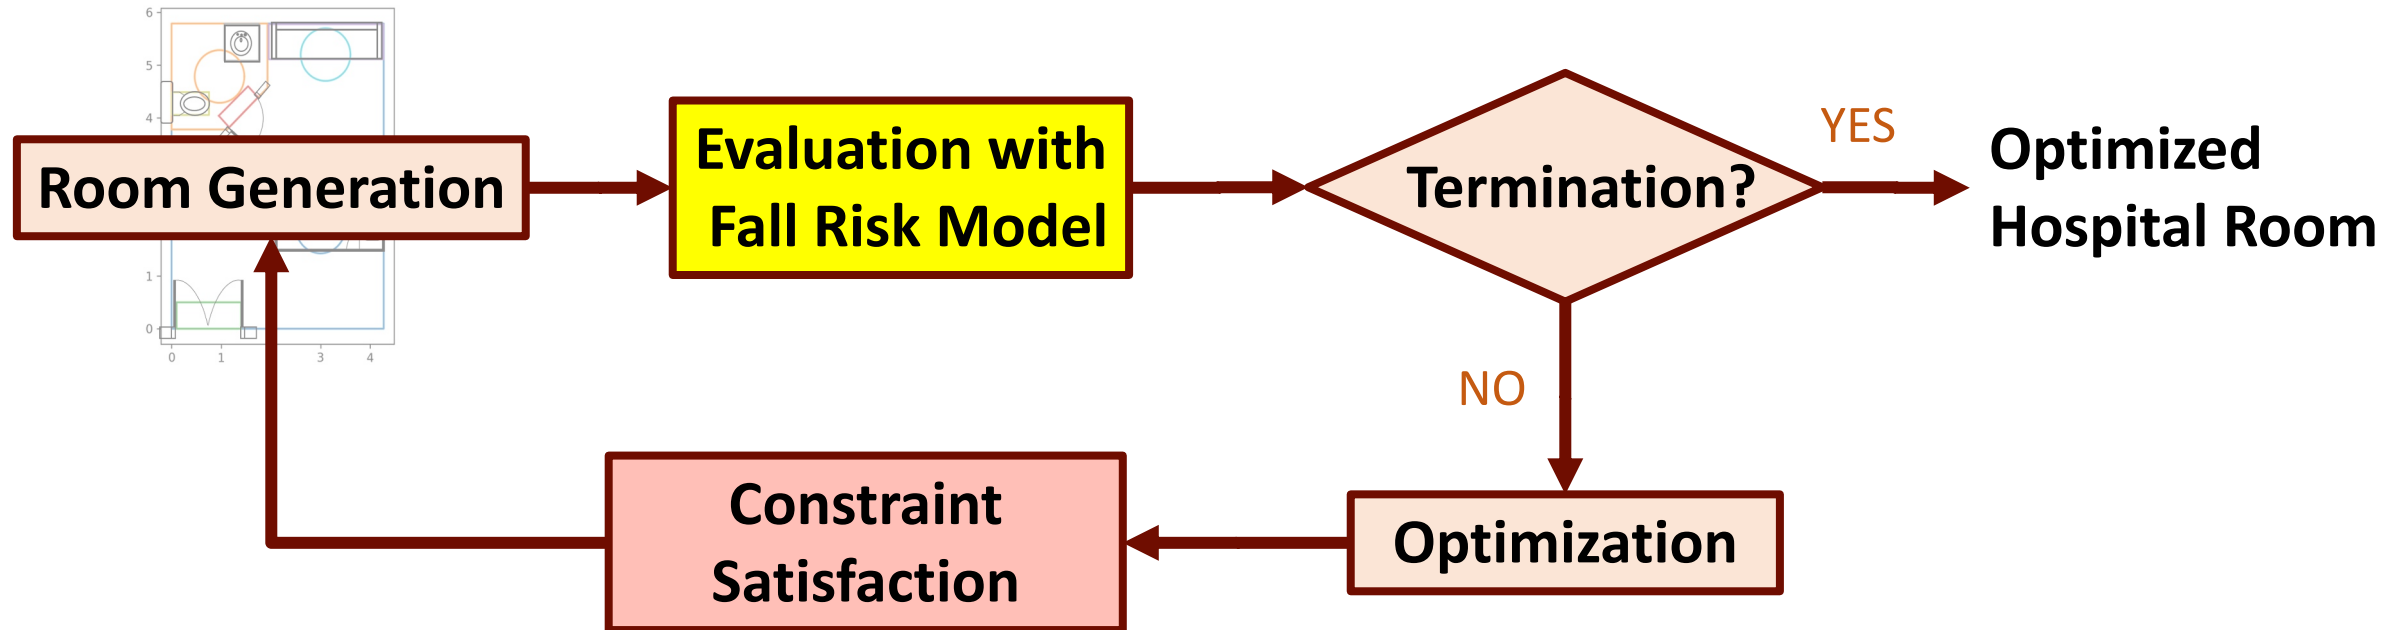
General Road Map





Initial Room




General Road Map



Fall Model (By: Roya Sabbaghnovin)




Baseline (Day)



Figures from: R. S. Novin, E. Taylor, T. Hermans, A. Merryweather, Development of a novel computational model for evaluating fall risk in patient room design., HERD: Health Environments Research & Design Journal (2020)

Fall Model: Baseline



Figures from: R. S. Novin, E. Taylor, T. Hermans, A. Merryweather, Development of a novel computational model for evaluating fall risk in patient room design., HERD: Health Environments Research & Design Journal (2020)

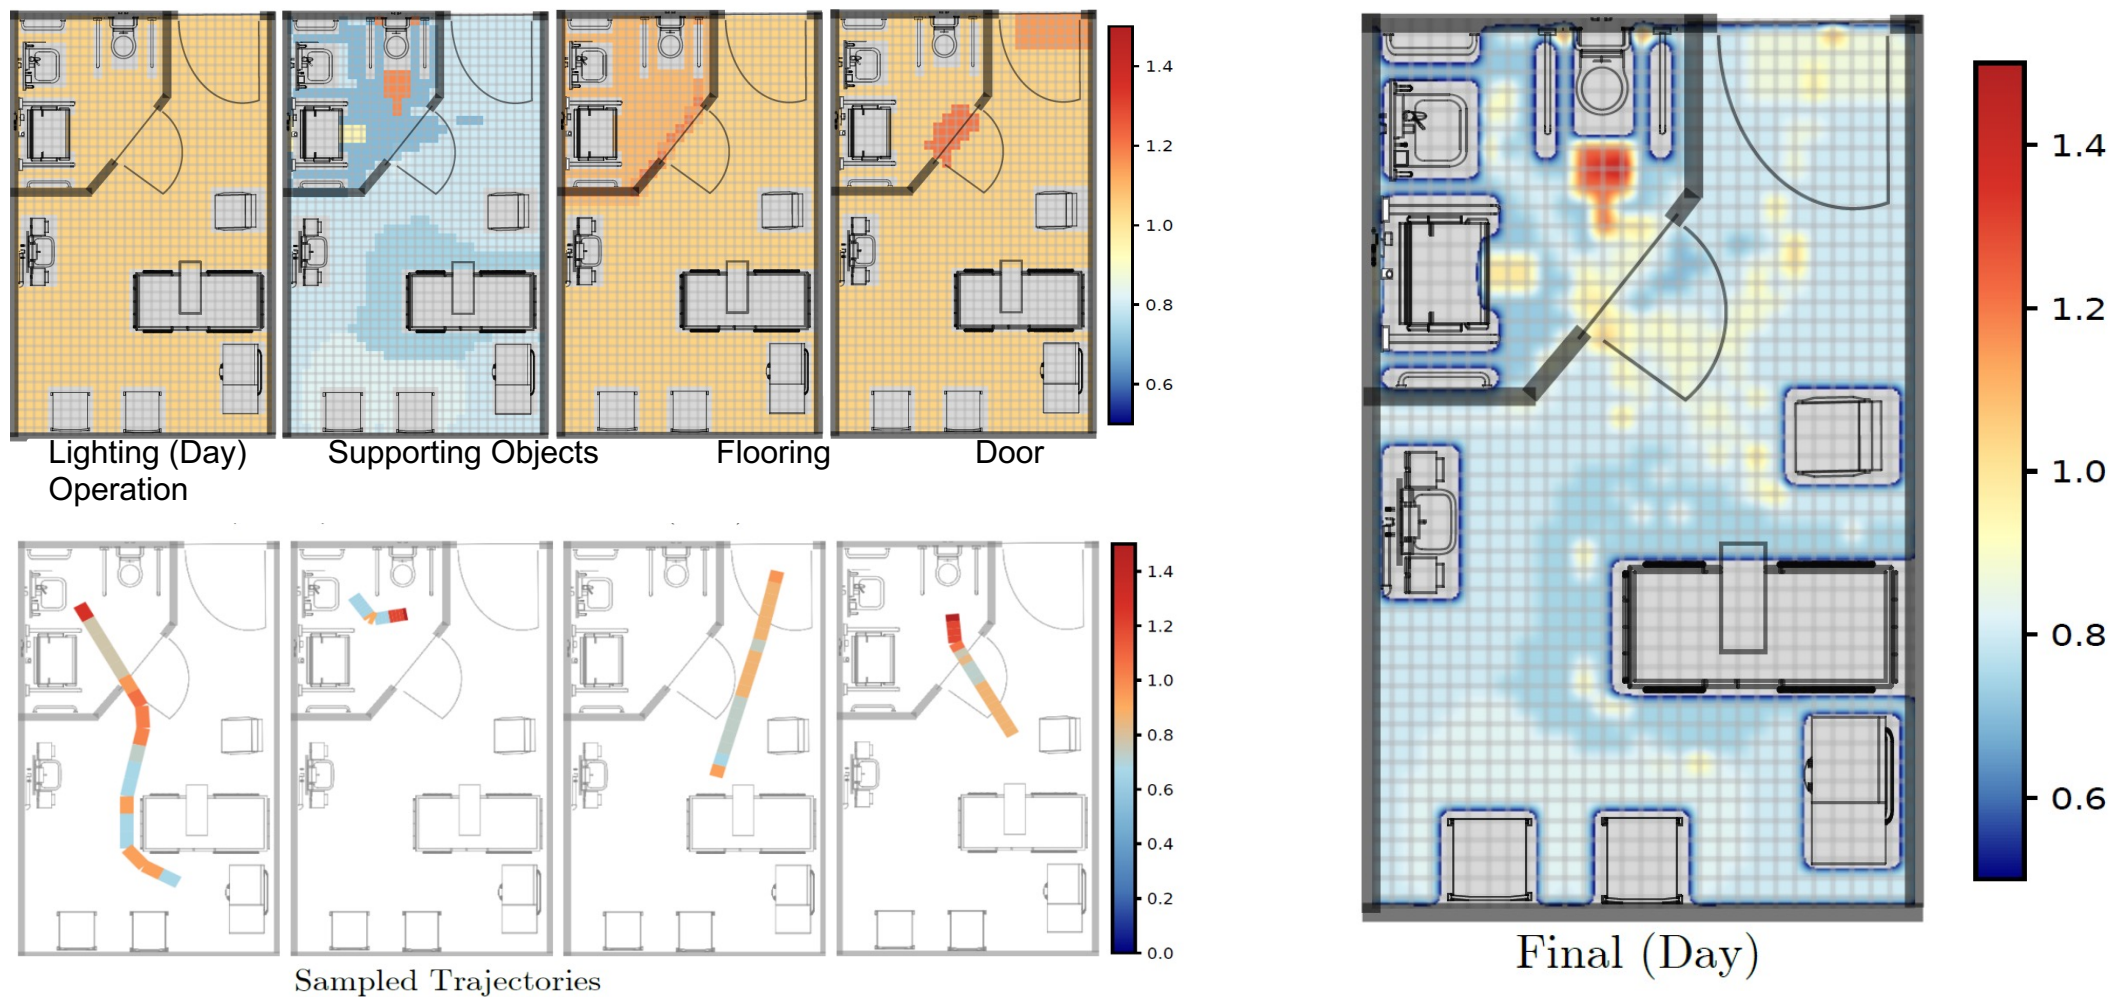
Fall Model: Trajectory



Sampled Trajectories

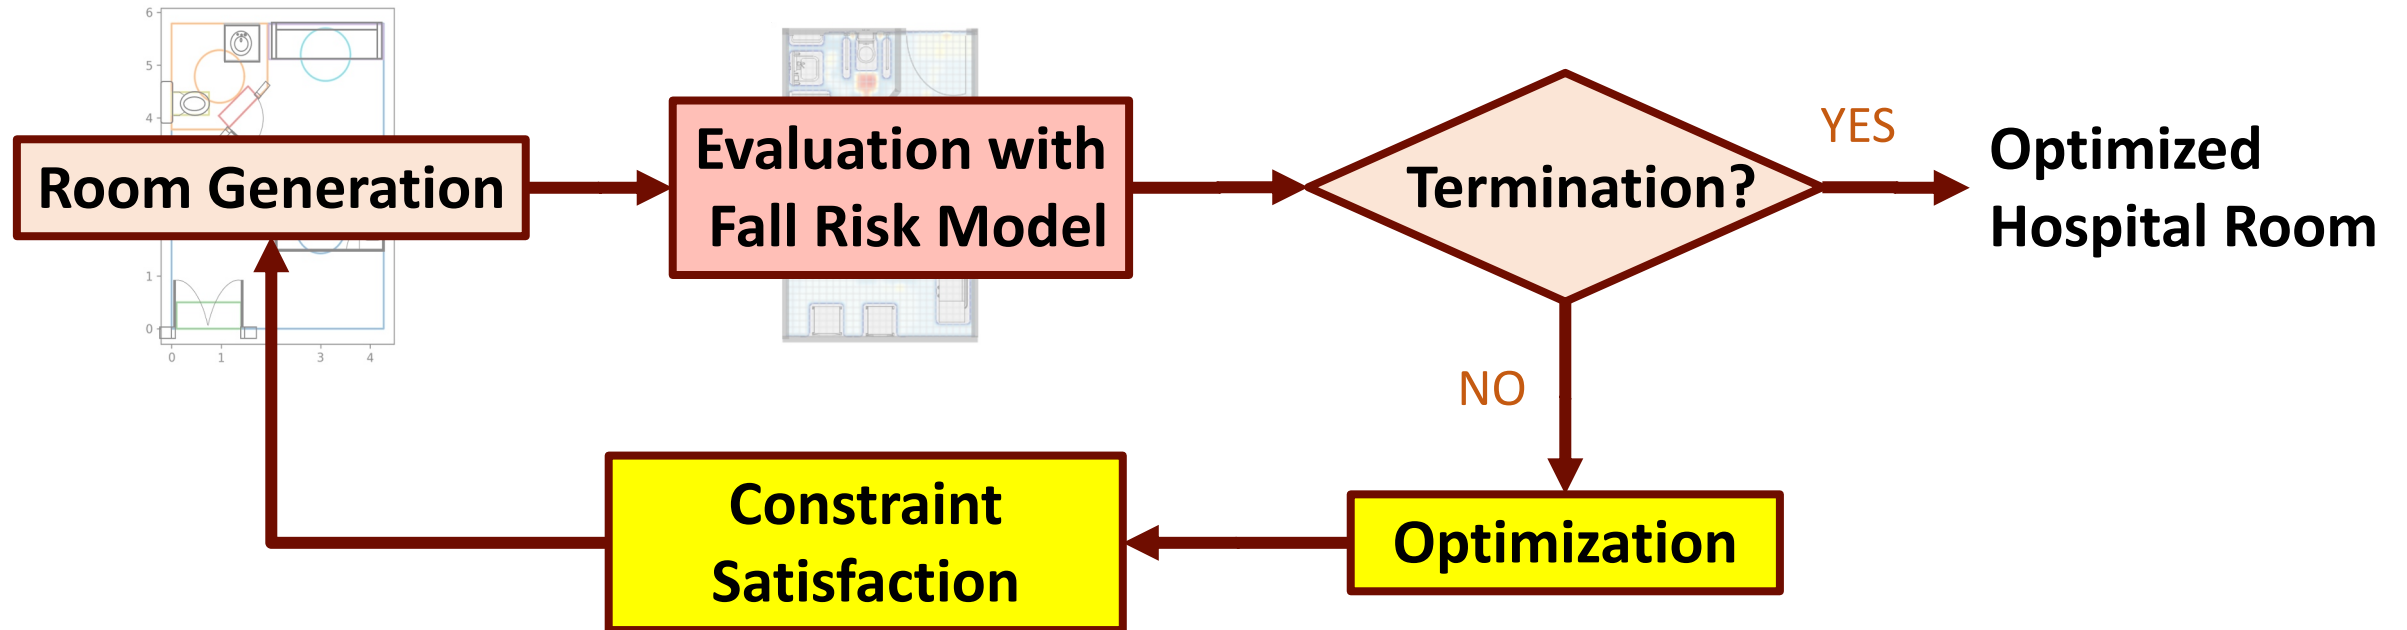
Figures from: R. S. Novin, E. Taylor, T. Hermans, A. Merryweather, Development of a novel computational model for evaluating fall risk in patient room design., HERD: Health Environments Research & Design Journal (2020)

Fall Model



Figures from: R. S. Novin, E. Taylor, T. Hermans, A. Merryweather, Development of a novel computational model for evaluating fall risk in patient room design., HERD: Health Environments Research & Design Journal (2020)

General Road Map



Constraint Satisfaction

- Set of variables: X
- Domain: D
- Set of Constraints: C

Constraint Satisfaction

- Set of variables: $X = \{F_0, \dots, F_n, L_0, \dots, L_m, D_{main}, D_{bath}\}$
- Domain: D
- Set of Constraints: C

Constraint Satisfaction


- Set of variables: $X = \{F_0, \dots, F_n, L_0, \dots, L_m, D_{main}, D_{bath}\}$
- Domain: D
 - $\overline{conf_{X_i}} = \{x, y, \theta\}$
 - $D_{X_i} = \left[\min(x_{room_{X_i}}), \max(x_{room_{X_i}}) \right], \left[\min(y_{room_{X_i}}), \max(y_{room_{X_i}}) \right], [0, 2\pi)$
- Set of Constraints: C

Constraint Satisfaction

- Set of variables: $X = \{F_0, \dots, F_n, L_0, \dots, L_m, D_{main}, D_{bath}\}$
- Domain: D
 - $\overline{conf_{X_i}} = \{x, y, \theta\}$
 - $D_{X_i} = \left[\min(x_{room_{X_i}}), \max(x_{room_{X_i}}) \right], \left[\min(y_{room_{X_i}}), \max(y_{room_{X_i}}) \right], [0, 2\pi]$
- Set of Constraints: C
 - Inside the room
 - Collision avoidance
 - Standing against a wall
 - Clearance
 - Door placement

Constraint Satisfaction


- Set of variables: $X = \{F_0, \dots, F_n, L_0, \dots, L_m, D_{main}, D_{bath}\}$
- Domain: D
 - $\overline{conf_{X_i}} = \{x, y, \theta\}$
 - $D_{X_i} = \left[\min(x_{room_{X_i}}), \max(x_{room_{X_i}}) \right], \left[\min(y_{room_{X_i}}), \max(y_{room_{X_i}}) \right], [0, 2\pi] \}$
- Set of Constraints: C
 - Inside the room
 - Collision avoidance
 - Standing against a wall
 - Clearance
 - Door placement



Constraint Satisfaction


Constraint Satisfaction Problem

- Set of variables: $X = \{F_0, \dots, F_n, L_0, \dots, L_m, D_{main}, D_{bath}\}$
- Domain: D
 - $\overline{conf_{X_i}} = \{x, y, \theta\}$
 - $D_{X_i} = \left[\min(x_{room_{X_i}}), \max(x_{room_{X_i}}) \right], \left[\min(y_{room_{X_i}}), \max(y_{room_{X_i}}) \right], [0, 2\pi] \}$
- Set of Constraints: C
 - Inside the room
 - Collision avoidance
 - Standing against a wall
 - Clearance
 - Door placement




Constraint Satisfaction

- Set of variables: $X = \{F_0, \dots, F_n, L_0, \dots, L_m, D_{main}, D_{bath}\}$
- Domain: D
 - $\overline{conf_{X_i}} = \{x, y, \theta\}$
 - $D_{X_i} = \left[\min(x_{room_{X_i}}), \max(x_{room_{X_i}}) \right], \left[\min(y_{room_{X_i}}), \max(y_{room_{X_i}}) \right], [0, 2\pi] \}$
- Set of Constraints: C
 - Inside the room
 - Collision avoidance
 - Standing against a wall
 - Clearance
 - Door placement




Constraint Satisfaction

- Set of variables: $X = \{F_0, \dots, F_n, L_0, \dots, L_m, D_{main}, D_{bath}\}$
- Domain: D
 - $\overline{conf_{X_i}} = \{x, y, \theta\}$
 - $D_{X_i} = [\min(x_{room_{X_i}}), \max(x_{room_{X_i}})], [\min(y_{room_{X_i}}), \max(y_{room_{X_i}})], [0, 2\pi]$
- Set of Constraints: C
 - Inside the room
 - Collision avoidance
 - Standing against a wall
 - Clearance
 - Door placement



Constraint Satisfaction

- Set of variables: $X = \{F_0, \dots, F_n, L_0, \dots, L_m, D_{main}, D_{bath}\}$
- Domain: D
 - $\overline{conf_{X_i}} = \{x, y, \theta\}$
 - $D_{X_i} = \left[\min(x_{room_{X_i}}), \max(x_{room_{X_i}}) \right], \left[\min(y_{room_{X_i}}), \max(y_{room_{X_i}}) \right], [0, 2\pi] \}$
- Set of Constraints: C
 - Inside the room
 - Collision avoidance
 - Standing against a wall
 - Clearance
 - Door placement



Constraint Satisfaction

Constraint Satisfaction Problem

- Set of variables: $X = \{F_0, \dots, F_n, L_0, \dots, L_m, D_{main}, D_{bath}\}$
- Domain: D
 - $\overline{conf_{X_i}} = \{x, y, \theta\}$
 - $D_{X_i} = \left[\min(x_{room_{X_i}}), \max(x_{room_{X_i}}) \right], \left[\min(y_{room_{X_i}}), \max(y_{room_{X_i}}) \right], [0, 2\pi] \}$
- Set of Constraints: C
 - Inside the room
 - Collision avoidance
 - Standing against a wall
 - Clearance
 - Door placement

What if the set of constraints cannot be satisfied?



Backtracking